

# Document Clustering

(Natural Language Processing)

HND Thilini  
hnd@ucsc.cmb.ac.lk



# Why document clustering?

- What happens when we want to categorize texts but have no data to learn from (i.e. no labelled/annotated, training data)
  - We relied on that when we used supervised learning for this task
- Can't we collect and annotate a dataset to learn from?
  - Easier said than done – an arduous task since we need lots of data to train on!
  - How much data? <http://machinelearningmastery.com/much--training--data--required--machine--learning>
- Clustering documents without 'supervision' of labels/annotations
  - Unsupervised learning algorithms

# How can we do document clustering?

- Can we measure similarity between documents?
- Can we use distance measures to find the most similar documents?
- Can we use such measures to group similar documents?
- Can we visualize such groups to determine their veracity?

# Key concepts

- Information retrieval

- Process of fetching the most relevant documents based on a query string
- This is what search engines do: relevancy is key

- Feature engineering

- Extracting abstractions or representations of documents from text content

- Similarity measures

- Ways to define and measure 'distance' between two words, phrases, sentences or documents

- Unsupervised machine learning

- Primarily clustering algorithms and dimensionality reduction algorithms

# Pipeline

- We use our normal preprocessing pipeline as before
- Normalizing
  - We add some words to our list of stopwords by looking at our domain
  - We also are only going to use text tokens using a regular expression – why?
- Feature extraction
  - We reuse our previous feature extractors but add the functionality provided by sklearn's Vectorizer classes
  - We include parameters for extracting features for n-grams instead of just unigrams (default) and setting max and min counts to be considered

# Text similarity

- Way to measure how similar two words, phrases, sentences etc are
- Two broad ways to define 'similarity'
  - Lexical level: the form of the relevant word/phrase/sentence
  - Semantic level: the meaning and context of the word/phrase/sentence
- Two main types of distance metrics used
  - Term similarity: measures distance between two words/tokens
  - Document similarity: measures the distance between entire documents
- We will be using multiple distance metrics and comparing results

# Term similarity

- Simple definition: how many characters match in sequence and in bag
  - i.e. a character vector and a bag of characters respectively
- We will use the most commonly used distance metrics to measure similarity
  - Hamming distance
  - Manhattan distance
  - Euclidean distance
  - Levenshtein (edit) distance
  - Cosine distance/similarity

# Hamming distance

- Number of positions that have different characters or symbols between *two strings of equal length*
  - Usually normalized for the length of the string

$$hd(u, v) = \sum_{i=1}^n (u_i \neq v_i)$$

$$norm\_hd(u, v) = \frac{\sum_{i=1}^n (u_i \neq v_i)}{n}$$





# Manhattan distance


- Distance between two points in a grid based on strictly horizontal or vertical paths
  - Not using the diagonal distance usually calculated by Euclidean distance
  - Also called city block distance, taxicab metric or L1 norm
- For strings, we subtract the difference between each pair of characters at each position of the two strings
  - So requires the strings to be of equal length
  - Can be normalized for length as before

$$md(u, v) = \|u - v\|_1 = \sum_{i=1}^n |u_i - v_i|$$

$$norm\_md(u, v) = \frac{\|u - v\|_1}{n} = \frac{\sum_{i=1}^n |u_i - v_i|}{n}$$

Manhattan Distance


$$|4 - 1| + |3 - 1| = 5$$



# Euclidean distance

- The shortest straight--line distance between two points
  - 'As the crow flies'
  - Also called Euclidean norm, L2 norm or L2 distance

$$ed(u, v) = \|u - v\|_2 = \sqrt{\sum_{i=1}^n (u_i - v_i)^2}$$



# Dealing with the 'length problem'

- Very unrealistic in general to assume this!
- Solution: 'edit' distance based measures
- Levenshtein distance – the most popular
  - The minimum number of edits needed (in the form of additions, deletions, or substitutions) to change/convert one term to the other
  - Importantly, the length of the two terms need not be the same
  - Minimum = difference in length between 2 terms
  - Maximum = length of the longer term
  - If equal, distance = 0
  - Satisfies 'triangle inequality'


P	A	R	T	Y
P	A	R	K	

Substitute

Delete

# Cosine distance

- A measure of the cosine of the angle between two terms when represented as non-zero positive vectors in an inner product space
  - The cosine similarity (cs) value will in general lie between -1 and +1
  - For bag-of-words (or characters) based models it will be between 0 and 1 since the frequencies can never be negative
  - Cosine distance is the inverse: i.e.  $1 - cs$



# Example of cosine similarity

- If  $d_1$  and  $d_2$  are two document vectors, then

$$\cos( d_1, d_2 ) = (d_1 \bullet d_2) / ||d_1|| ||d_2|| ,$$

where  $\bullet$  indicates vector dot product and  $||d||$  is the length of vector  $d$ .

- Example:

$$d_1 = \mathbf{3\ 2\ 0\ 5\ 0\ 0\ 0\ 2\ 0\ 0}$$

$$d_2 = \mathbf{1\ 0\ 0\ 0\ 0\ 0\ 0\ 1\ 0\ 2}$$

$$d_1 \bullet d_2 = 3*1 + 2*0 + 0*0 + 5*0 + 0*0 + 0*0 + 0*0 + 2*1 + 0*0 + 0*2 = 5$$

$$||d_1|| = (3*3 + 2*2 + 0*0 + 5*5 + 0*0 + 0*0 + 0*0 + 2*2 + 0*0 + 0*0)^{0.5} = (42)^{0.5} = 6.481$$

$$||d_2|| = (1*1 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 1*1 + 0*0 + 2*2)^{0.5} = (6)^{0.5} = 2.245$$

$$\cos( d_1, d_2 ) = .3150$$

# Document similarity

- We do this by vectorizing our documents using the `utils` module
  - Instead of vectorizing words in terms of characters
- Then we can use cosine similarity (not distance) to measure the similarity between two documents
  - Since we use bag-of-words models, the sentences need not be in the same word order (desirable feature)
- We will also use two other document similarity measures
  - Hellinger--Bhattacharya distance (HB--distance, Bhattacharya distance)
  - Okapi BM25 ranking (BM25)

# Document similarity

- For real-world problems dealing with large documents there are many optimized metrics that can (and should) be used
  - `cosine_similarity()` function of `sklearn.metrics.pairwise`
  - `gensim`'s `similarities` module or `cossim()` function from `gensim.matutils` module
  - `gensim`'s `hellinger()` function from the `gensim.matutils` module
  - `gensim.summarization` package has an implementation of `bm25`
- Try loading bigger corpora and test and compare the output of these functions
- Used in IR systems – e.g. in `solr`, `Elasticsearch` (on top of `Lucene`)
  - See <https://www.elastic.co/blog/found-bm-vs-lucene-default-similarity>

# Document clustering


- In text/document classification, we categorize text/documents using pre-labelled ‘training data’
- What happens if we don’t have the luxury of annotated data?
- That is where document clustering is needed
  - We need unsupervised learning to solve this
- There are multiple ways to cluster any collection of data
- Many clustering algorithms usually require you to specify the number of clusters – why may this be not so bad?
- Evaluating how good a clustering scheme automatically is hard




# Ambiguity of clustering




How many clusters?




Two Clusters



Four Clusters





Six Clusters



# Ambiguity of clustering – paritional/hard








Input Data



A Partitional Clustering



# Ambiguity of clustering – hierarchical



Clustering Solution 2


# Clustering algorithms

- Hierarchical clustering
  - Visualized using a *dendrogram*




# Clustering algorithms

- Centroid based clustering
  - Build clusters using central representative members per cluster (k-means)




# Clustering algorithms

- Distribution based clustering
  - Uses the concept that objects having similar distributions should be together



# Clustering algorithms

- Density based clustering
  - Uses the high density areas in the space for clustering (DBSCAN)



# Clustering algorithms

- Those interested can refer to following links for general info
  - [http://scikit--learn.org/stable/auto\\_examples/cluster/plot\\_cluster\\_comparison.html](http://scikit--learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html)
  - [http://hdbscan.readthedocs.io/en/latest/comparing\\_clustering\\_algorithms.html](http://hdbscan.readthedocs.io/en/latest/comparing_clustering_algorithms.html)
  - <https://www.toptal.com/machine--learning/clustering--algorithms>
- You are encouraged to try some of the other algorithms available via scikit--learn



# K-means clustering


- Objective - group similar data points together and discover underlying patterns
- Tries to minimize within-cluster sum-of-squares (inertia)
- As its name implies, the number of clusters ( $k$ ) needs to be specified
- This is so with all centroid--based algorithms
- Popular because it is highly scalable

# K-means clustering

- For a dataset with  $N$  points
  - **Step-1:** Select the number  $K$  to decide the number of clusters.
  - **Step-2:** Select random  $K$  points or centroids.
  - **Step-3:** Assign each data point to their closest centroid, which will form the predefined  $K$  clusters.
  - **Step-4:** Calculate the variance and place a new centroid of each cluster.
  - **Step-5:** Repeat the third steps, which means reassign each datapoint to the new closest centroid of each cluster.
  - **Step-6:** If any reassignment occurs, then go to step-4 else go to FINISH.
  - **Step-7:** The model is ready.

# K-means clustering


- Suppose we have two variables M1 and M2.



Suppose  $k = 2$



# K-means clustering

- Choose 2 points (k points) or centroid to form the cluster.
- These points can be either the points from the dataset or any other point




# K-means clustering

- Assign each data point of the scatter plot to its closest K-point or centroid.
- Compute it by calculating the distance between two points.
- Draw a median between both the centroids




# K-means clustering

- Find the closest cluster - repeat the process by choosing a **new centroid**




# K-means clustering

- Reassign each datapoint to the new centroid.



# K-means clustering


- Reassign each datapoint to the new centroid.






# K-means clustering

- Repeat the process



# K-means clustering

- Model is formed



# K-means algorithm (contd.)

- Can converge to a 'local optimum'
- In practice, we run this algorithm multiple times with different epochs
- As common, convergence is highly dependent on initial assignment
- Two common solutions
  - Take average of multiple iterations with multiple random initializations
  - Initialize the centroids to be far apart from each other (kmeans++)
    - Implemented in scikit-learn
- Cannot visualize owing to the data generally being high-dimensional
  - We use PCA or MDS (multidimensional scaling) to view in 2 or 3 dimensions

# Choosing the value of K

- Performance of the K-means clustering algorithm depends upon highly efficient clusters that it forms
- Choosing the optimal number of clusters is a big task
- Some methods that can be used
  - The Elbow method
  - The Silhouette Method
  - ...

# The Elbow Method


- Probably the most well-known method for determining the optimal number of clusters
- Calculate the **Within Cluster Sum of Squares (WCSS)**

$$WCSS = \sum_{P_i \text{ in Cluster } 1} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster } 2} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster } 3} \text{distance}(P_i, C_3)^2$$

- $\sum_{P_i \text{ in Cluster } 1} \text{distance}(P_i, C_1)^2$ : is the sum of the square of the distances between each data point and its centroid within a cluster1
- To measure the distance between data points and centroid  
Euclidean distance, Manhattan distance, etc.

# The Elbow Method

- Elbow method follows the below steps:
- It executes the K-means clustering on a given dataset for different K values (ranges from 1-10).
- For each value of K, calculates the WCSS value.
- Plots a curve between calculated WCSS values and the number of clusters K.
- The sharp point of bend or a point of the plot looks like an arm, then that point is considered as the best value of K.



# Affinity Propagation algorithm

- AP tries to build clusters without a pre-specified number
  - Uses the concept of 'message passing' between data points
- Finding *exemplars* that act as representatives of dataset
  - Based on the messages passed between each pair of data point
  - Hence, convergence takes a long time
- We use cosine similarity as the measure in our message passing
- Related to *spectral clustering*

# Affinity Propagation (contd.)

- Iteratively proceed by executing the following message passing steps:
  - Send 'responsibility' updates between pairs of data points as a matrix R
  - Send 'availability' updates by as a matrix A to all data pairs
  - Keep doing these steps until convergence

$$r(i, k) \leftarrow \text{sim}(i, k) - \max_{k' \neq k} \{a(i, k') + \text{sim}(i, k')\}$$


$$a(i, k) \leftarrow \min \left( 0, r(k, k) + \sum_{i' \in [i, k]} \max(0, r(i', k)) \right) \text{ for } i \neq k$$

- Task: find the best intuitive description of the AP algorithm online!



# Ward's Agglomerative Hierarchical Clustering


- Hierarchical clustering attempts to build nested hierarchies
- Two broad types:
  - Agglomerative (bottom--up): each point in its own class initially
  - Divisive (top--down): all points in one class initially
- Merges and splits (respectively) happen in 'greedy' fashion



# Ward's Agglomerative Hierarchical Clustering

- Deciding which points (or clusters) should join:
  - Based on a *distance metric* to measure the similarity or dissimilarity degree between data points (we use cosine similarity)
  - Using a *linkage criterion* that determines the metric to be used for the merging strategy of clusters (we use Ward's method)
- Ward's method minimized the sum of squared differences within all the clusters (variance minimization)
  - Minimize variance in each cluster using objective function such as L2 norm
- We use `matplotlib` to plot the final resulting dendrogram

# Ward's Agglomerative Hierarchical Clustering



# Cluster validation

- Comparing the clustering results to *ground truth* (externally known results)
  - External Index
- Evaluating the quality of clusters *without* reference to external information
  - Use only the data
  - Internal Index

# Cluster validation – external

- **Notation**
  - $N$ : number of objects in the data set
  - $P=\{P_1,\dots,P_s\}$ : the set of “ground truth” clusters
  - $C=\{C_1,\dots,C_t\}$ : the set of clusters reported by a clustering algorithm
- **The “incidence matrix”**
  - $N \times N$  (both rows and columns correspond to objects)
  - $P_{ij} = 1$  if  $O_i$  and  $O_j$  belong to the same “ground truth” cluster in  $P$ ;  $P_{ij}=0$  otherwise
  - $C_{ij} = 1$  if  $O_i$  and  $O_j$  belong to the same cluster in  $C$ ;  $C_{ij}=0$  otherwise

# Cluster validation – external (contd.)

- A pair of data object  $(O_i, O_j)$  falls into one of the following categories

- SS:  $C_{ij}=1$  and  $P_{ij}=1$ ; (agree)
- DD:  $C_{ij}=0$  and  $P_{ij}=0$ ; (agree)
- SD:  $C_{ij}=1$  and  $P_{ij}=0$ ; (disagree)
- DS:  $C_{ij}=0$  and  $P_{ij}=1$ ; (disagree)

- **Rand index** 
$$Rand = \frac{|Agree|}{|Agree| + |Disagree|} = \frac{|SS| + |DD|}{|SS| + |SD| + |DS| + |DD|}$$

- may be dominated by DD

- **Jaccard Coefficient** 
$$Jaccard\ coefficient = \frac{|SS|}{|SS| + |SD| + |DS|}$$

# Cluster validation – external (contd.)

Clustering

	g 1	g 2	g 3	g 4	g 5
g 1	1	1	1	0	0
g 2	1	1	1	0	0
g 3	1	1	1	0	0
g 4	0	0	0	1	1
g 5	0	0	0	1	1

Ground truth



Groundtruth

	g 1	g 2	g 3	g 4	g 5
g 1	1	1	0	0	0
g 2	1	1	0	0	0
g 3	0	0	1	1	1
g 4	0	0	1	1	1
g 5	0	0	1	1	1

Clustering

	Same Cluster	Different Cluster
Same Cluster	9	4
Different Cluster	4	8

$$Rand = \frac{|SS| + |DD|}{|SS| + |SD| + |DS| + |DD|} = \frac{17}{25}$$

$$Jaccard = \frac{|SS|}{|SS| + |SD| + |DS|} = \frac{9}{17}$$

# Cluster validation – internal

- **Cohesion** is measured by the within cluster sum of squares

$$WSS = \sum_i \sum_{x \in C_i} (x - m_i)^2$$

- **Separation** is measured by the between cluster sum of squares


$$BSS = \sum_i |C_i| (m - m_i)^2$$

where  $|C_i|$  is the size of cluster  $i$ ,  $m$  is the centroid of the whole data set

- $BSS + WSS = \text{constant}$
- WSS (Cohesion) measure is called Sum of Squared Error (SSE)—a commonly used measure
- A larger number of clusters tend to result in smaller SSE



# Cluster validation – internal (contd.)



**K=1 :**

$$WSS = (1-3)^2 + (2-3)^2 + (4-3)^2 + (5-3)^2 = 10$$

$$BSS = 4 \times (3-3)^2 = 0$$

$$Total = 10 + 0 = 10$$

**K=2 :**

$$WSS = (1-1.5)^2 + (2-1.5)^2 + (4-4.5)^2 + (5-4.5)^2 = 1$$

$$BSS = 2 \times (3-1.5)^2 + 2 \times (4.5-3)^2 = 9$$

$$Total = 1 + 9 = 10$$

**K=4:**

$$WSS = (1-1)^2 + (2-2)^2 + (4-4)^2 + (5-5)^2 = 0$$

$$BSS = 1 \times (1-3)^2 + 1 \times (2-3)^2 + 1 \times (4-3)^2 + 1 \times (5-3)^2 = 10$$


$$Total = 0 + 10 = 10$$

# Cluster validation – internal (contd.)

- Silhouette Coefficient combines ideas of both cohesion and separation
- For an individual point,  $i$ 
  - Calculate  $a$  = average distance of  $i$  to the points in its cluster
  - Calculate  $b$  = min (average distance of  $i$  to points in another cluster)
  - The **silhouette coefficient** for a point is then given by

$$s = 1 - a/b \quad \text{if } a < b, \quad (s = b/a - 1 \quad \text{if } a \geq b, \text{ not the usual case})$$


- Typically between 0 and 1
- The closer to 1 the better




- Can calculate the Average Silhouette width for a cluster or a clustering

# Are there clusters in the data?


Random  
Points




DBSCAN



K-means



Complete  
Link



# Application: greatest movies of all time!

- IMDb contains details about movies and TV series
  - Together with reviews and ratings
- We will work with the 'ultimate list' (100 greatest movies)
  - <http://www.imdb.com/list/ls055592025/>
- We want to cluster the movies based on the title and synopsis
  - Already cleaned by various others (available as movie\_data.csv)
- We load, normalize and extract features as before
  - Using bigram TF-IDF weights with min and max cutoffs
  - Then we use our clustering algorithms and compare results